

# Data Engineering 2020 Fall

LTAT.02.007

Ass Prof. Riccardo Tommasini

Assistants: **Fabiano Spiga, Mohamed Ragab, Hassan Eldeeb**



[https://courses.cs.ut.ee/2020/  
dataeng](https://courses.cs.ut.ee/2020/dataeng)

Forum

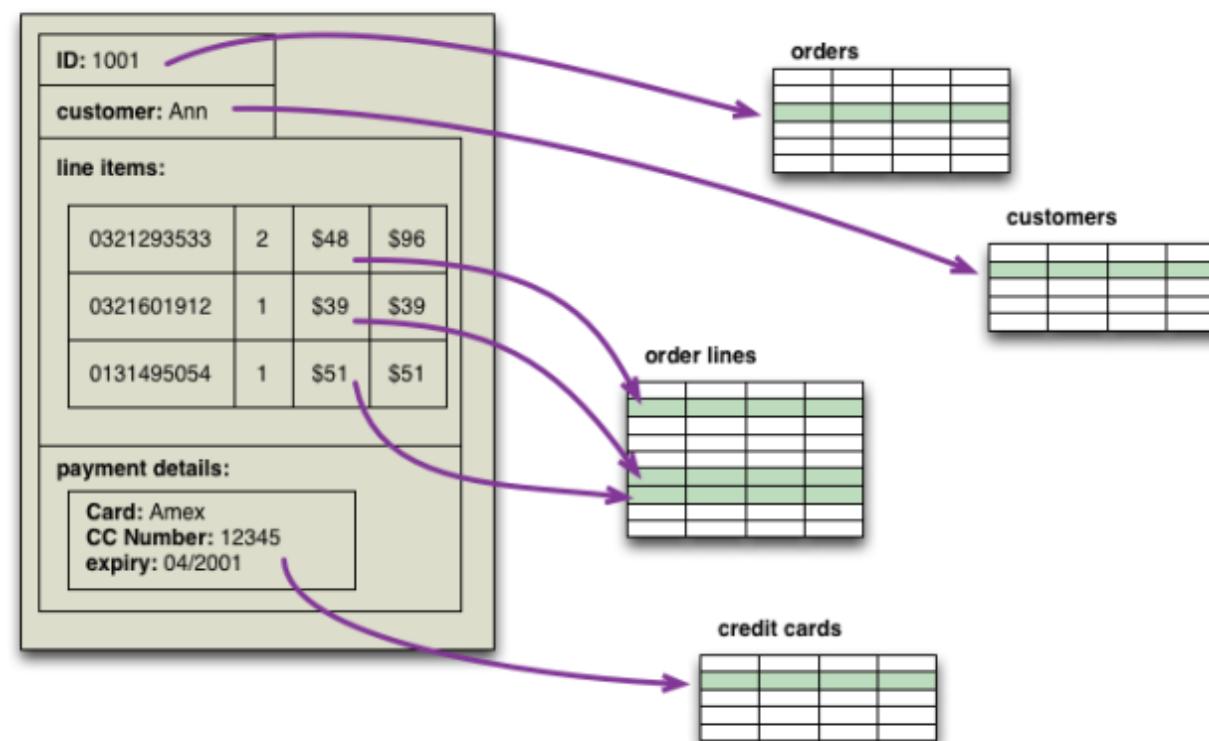
Moodle



# Document Databases

# Why document-based?

- Handles Schema Changes Well (easy development)
- Solves Impedance Mismatch problem
- Rise of JSON
- python module: simplejson



# What is a document?

```
{  
  "business_id": "rncjoVoEFUJGUoC1JgnUA",  
  "full_address": "8466 W Peoria AvenSte 6nPeoria, AZ 85345",  
  "open": true,  
  "categories": ["Accountants", "Professional Services", "Tax Services"],  
  "city": "Peoria",  
  "review_count": 3,  
  "name": "Peoria Income Tax Service",  
  "neighborhoods": [],  
  "longitude": -112.241596,  
  "state": "AZ",  
  "stars": 5.0,  
  "latitude": 33.58186700000003,  
  "type": "business":  
}
```

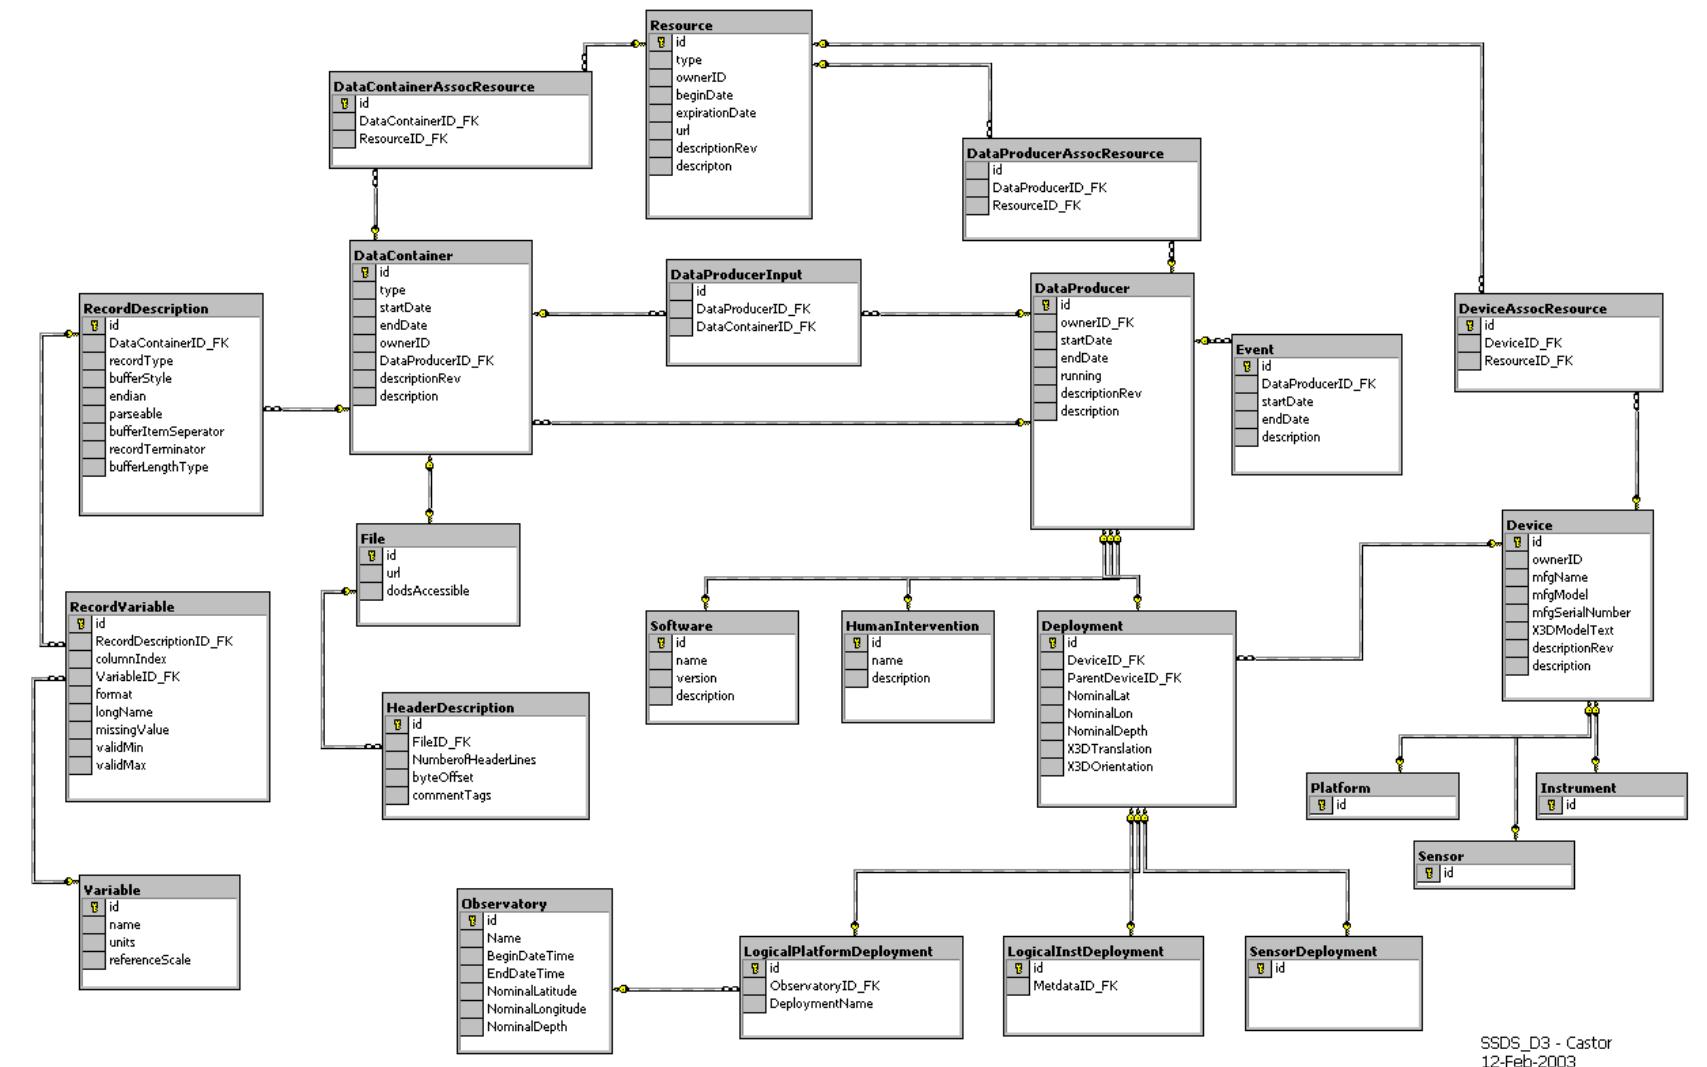
# JSON Format

- Data is in name / value pairs
- A name/value pair consists of a field name followed by a colon, followed by a value:
  - Example: "name": "R2-D2"
- Data is separated by commas
  - Example: "name": "R2-D2", race : "Droid"
- Curly braces hold objects
  - Example: {"name": "R2-D2", race : "Droid", affiliation: "rebels"}
- An array is stored in brackets []
  - Example [ {"name": "R2-D2", race : "Droid", affiliation: "rebels"}, {"name": "Yoda", affiliation: "rebels"} ]

# Designing NoSQL Data Structures

- NoSQL data structures driven by application design.
  - Need to take into account necessary CRUD operations
- To embed or not to imbed. That is the question!
  - Rule of thumb is to imbed whenever possible.
  - No modeling standards or CASE tools!

# Relational to Document



```
{
  "title" : "MongoDB",
  "contributors": [
    { "name" : "Eliot Horowitz",
      "email" : "eliot@10gen.com" },
    { "name" : "Dwight Merriman"
      "email" : "dwight@10gen.com" } ],
  "model" : {
    "relational" : false,
    "awesome" : true
  }
}
```

# A normalized structure

```
{  
  "_id"      : "First Post",  
  "author"   : "Rick",  
  "text"     : "This is my first post."  
  
}  
  
{  
  "_id"    : ObjectId(...),  
  "post_id": "First Post",  
  "author" : "Bob",  
  "text"   : "Nice Post!"  
}
```

# A (denormalized) embedded structure

```
{  
  
  "_id" : "First Post",  
  "comments" : [  
    { "author" : "Bob",  
      "text" : "Nice Post!"},  
    { "author" : "Tom",  
      "text" : "Dislike!"}],  
  "comment_count" : 2  
}
```

# A polymorphic structure

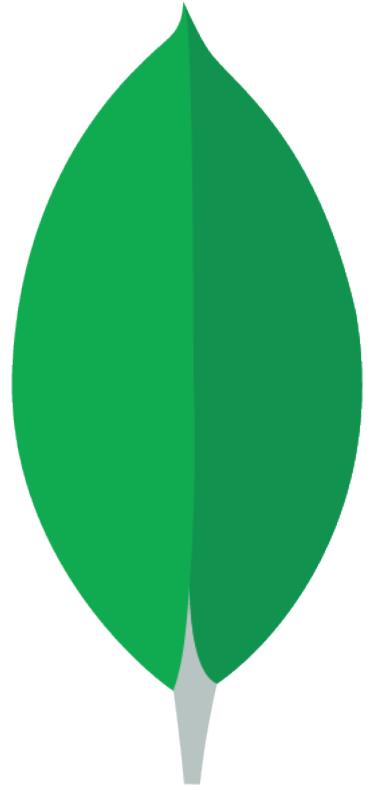
- When all the documents in a collection are similarly, but not identically structured.
- Enables simpler schema migration.
- no more of custom \_field\_1
- Better mapping of object - oriented inheritance and polymorphism.

```
{  
    "_id" : 1,  
    "title": "Welcome",  
    "url": "/",  
    "type": "page",  
    "content": "Welcome to my wonderful wiki."  
}  
  
{  
    "_id": 3,  
    "title": "Cool Photo",  
    "url": "/photo.jpg",  
    "type": "photo",  
    "content": Binary(...)  
}
```

# List of Systems

- **MongoDB**
- CouchDB
- OrientDB

# MongoDB



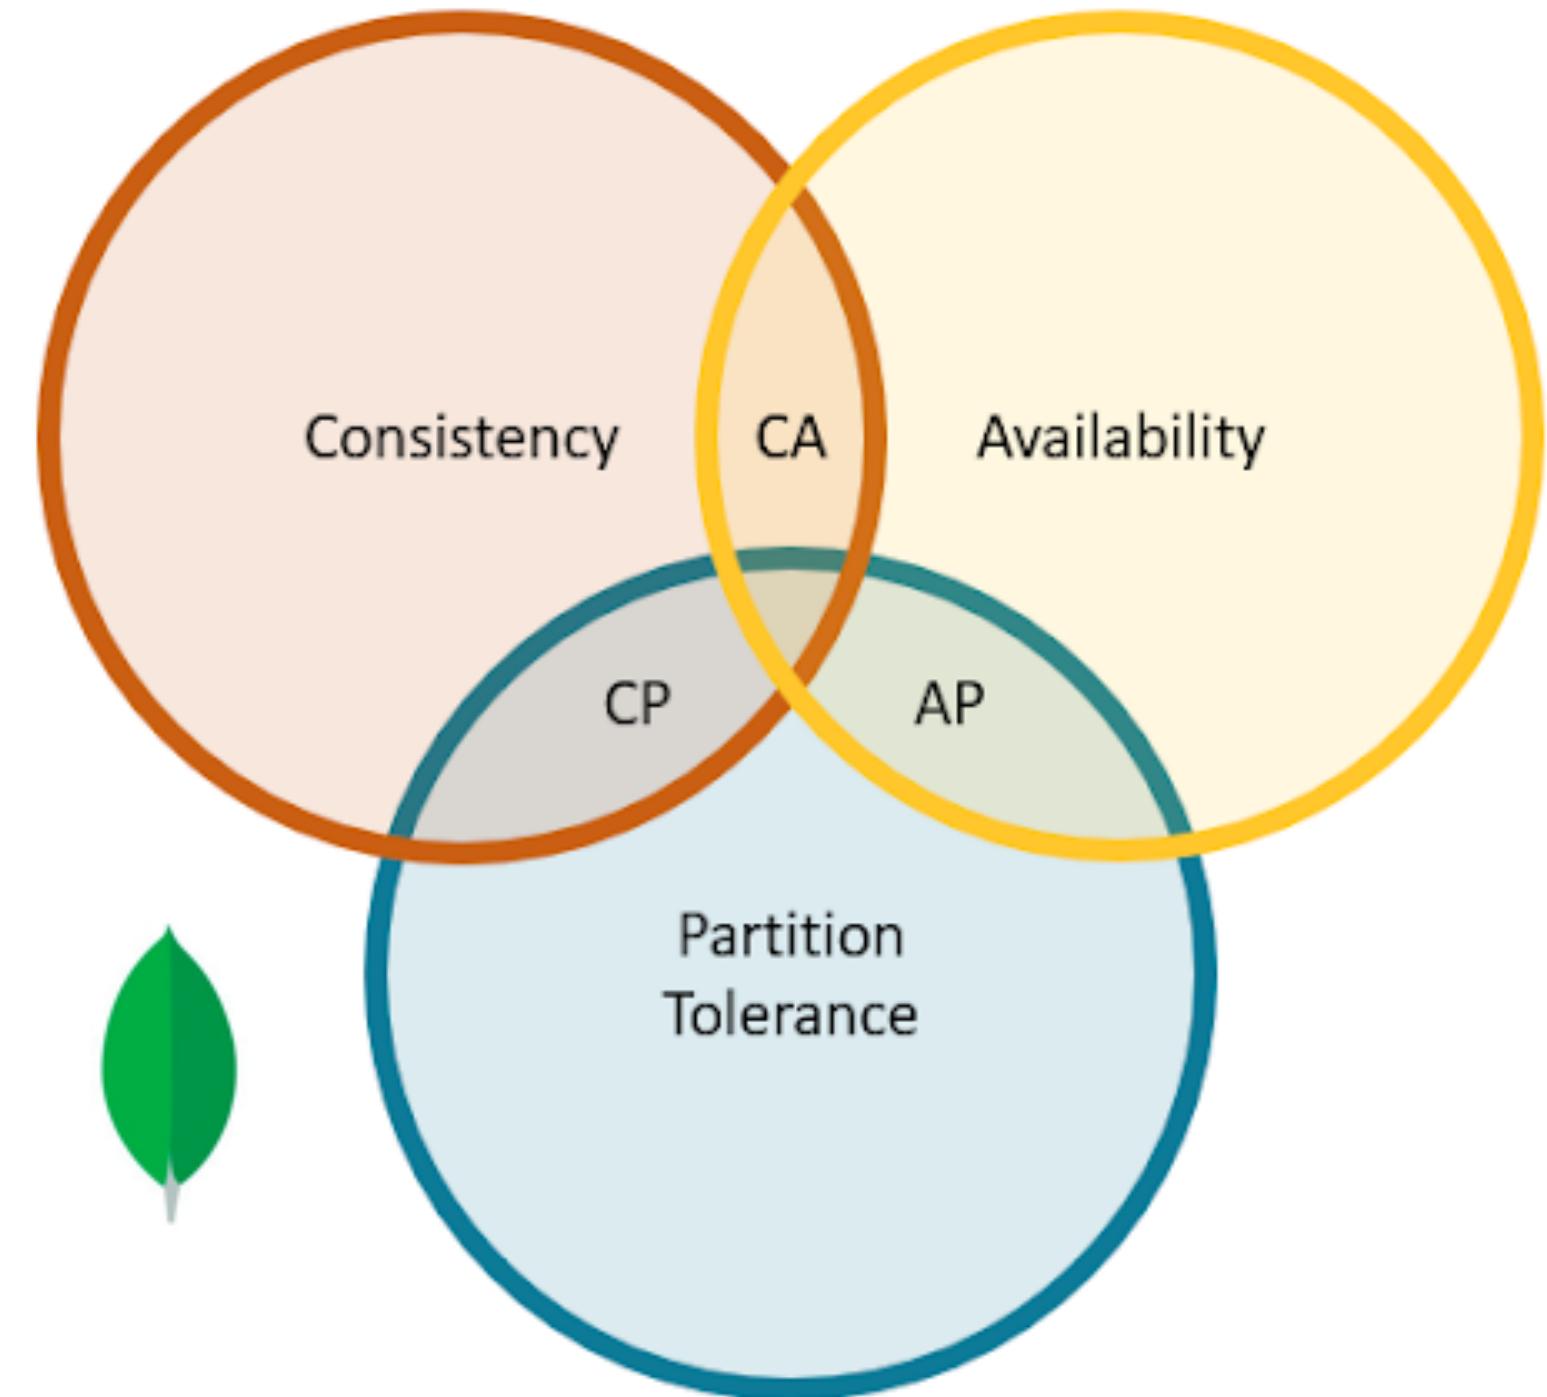
# mongoDB®

# History and Motivation

- An open source and document-oriented database.
- Data is stored in JSON-like documents.
- Designed with both scalability and developer agility.
- Dynamic schemas.
- Automatic data sharding

# What MongoDB is :

- An In-Memory [[Document Databases]]
- Strong consistency (**C**)
- *Tuneably* available (**A**)
- Horizontal Scalable (**P**)



# What MongoDB is not

- Always Available<sup>91</sup>
- No Schemas
- No transactions
- No joins
- Max document size of 16MB<sup>92</sup>

---

<sup>91</sup> Larger documents handled with GridFS

<sup>92</sup> there will always be downtime when (i) the new leader is getting elected or (ii) the client driver disconnects from the leader

# Use Cases

-  Capture **game** events, scaling to meet high-write workloads.
-  Financial Services: Risk Analytics & Reporting, Trade Repository
-  **BOSCH** manufacturing, automotive, retail, and energy
- **ThermoFisher** SCIENTIFIC fast-changing sensor data captured from multiple devices and experiments

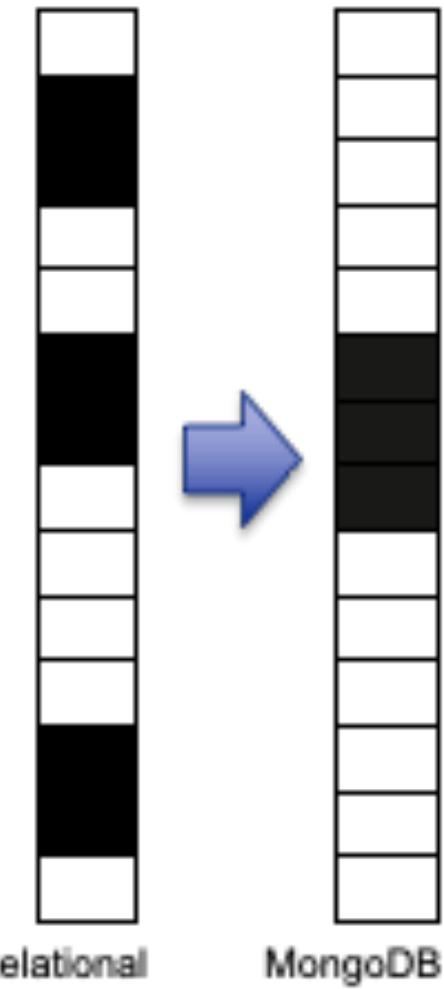
# When to consider MongoDB

- When you need high availability of data
- when you need fast and instant data recovery
- when do not want to sustain schema migration costs

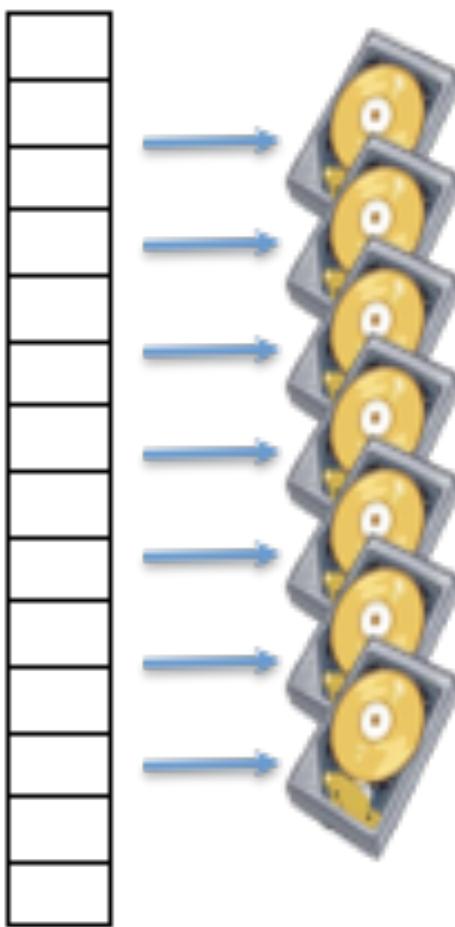
# Advantages

- Full featured indexes
- Sophisticated query language
- Easy mapping to object oriented code
- Native language drivers in all popular language
- Simple to setup and manage
- Operates at in-memory speed wherever possible
- Auto-sharding built in
- Dynamically add / remove capacity with no downtime

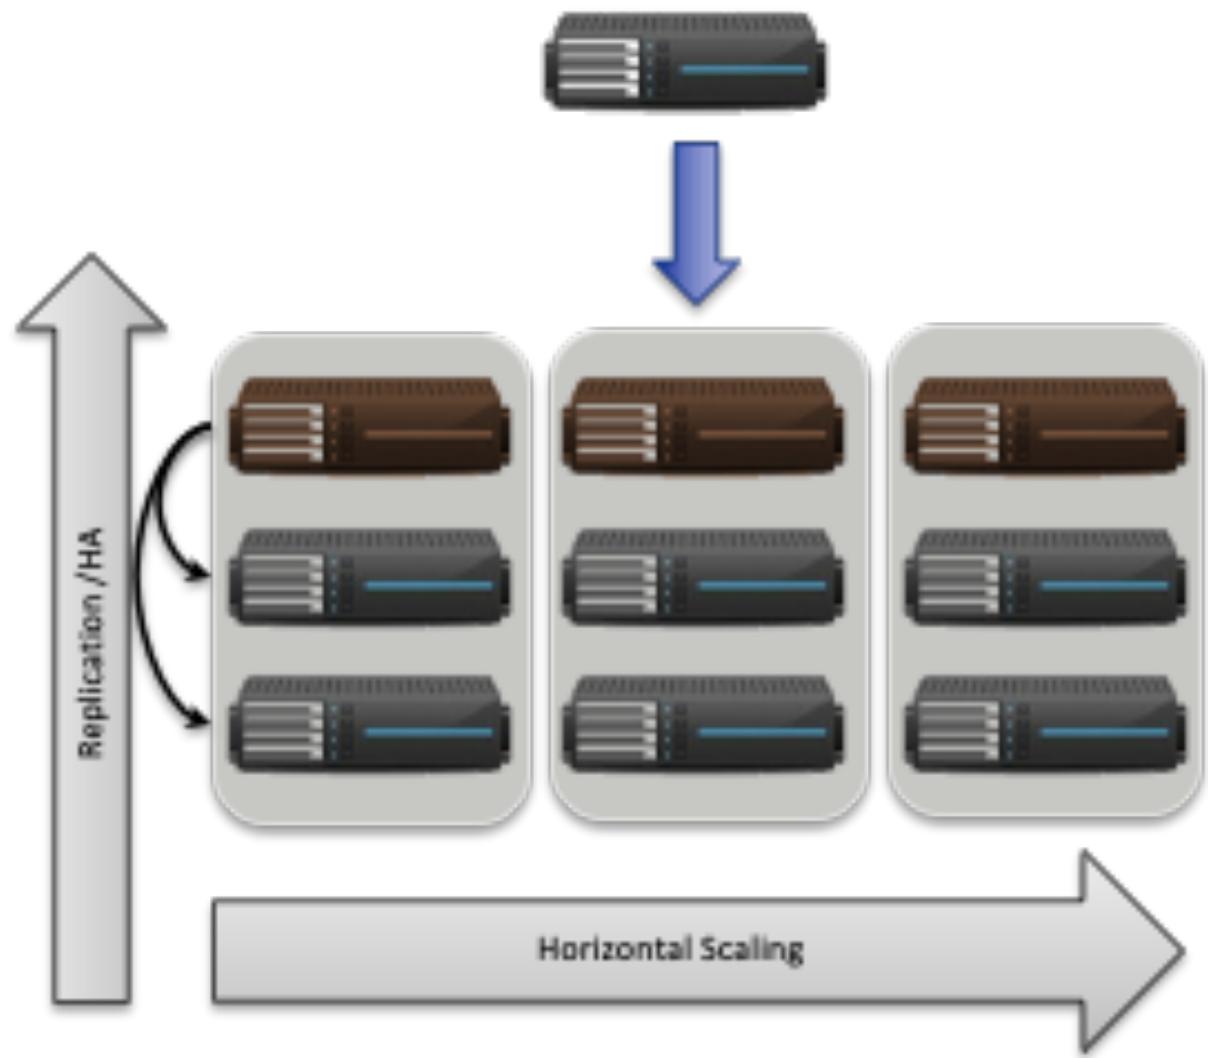
Better data locality



In-Memory  
Caching



Distributed Architecture



# Terminology: SQL vs MongoDB

## SQL Terms/Concepts

database

table

row

column

index

table joins (e.g. select queries)

Primary keys

Aggregation (e.g. group by)

## MongoDB Terms/Concepts

database

collection

document

field

index

embedded documents and linking

\_id field is always the primary key

aggregation pipeline

# Data Model

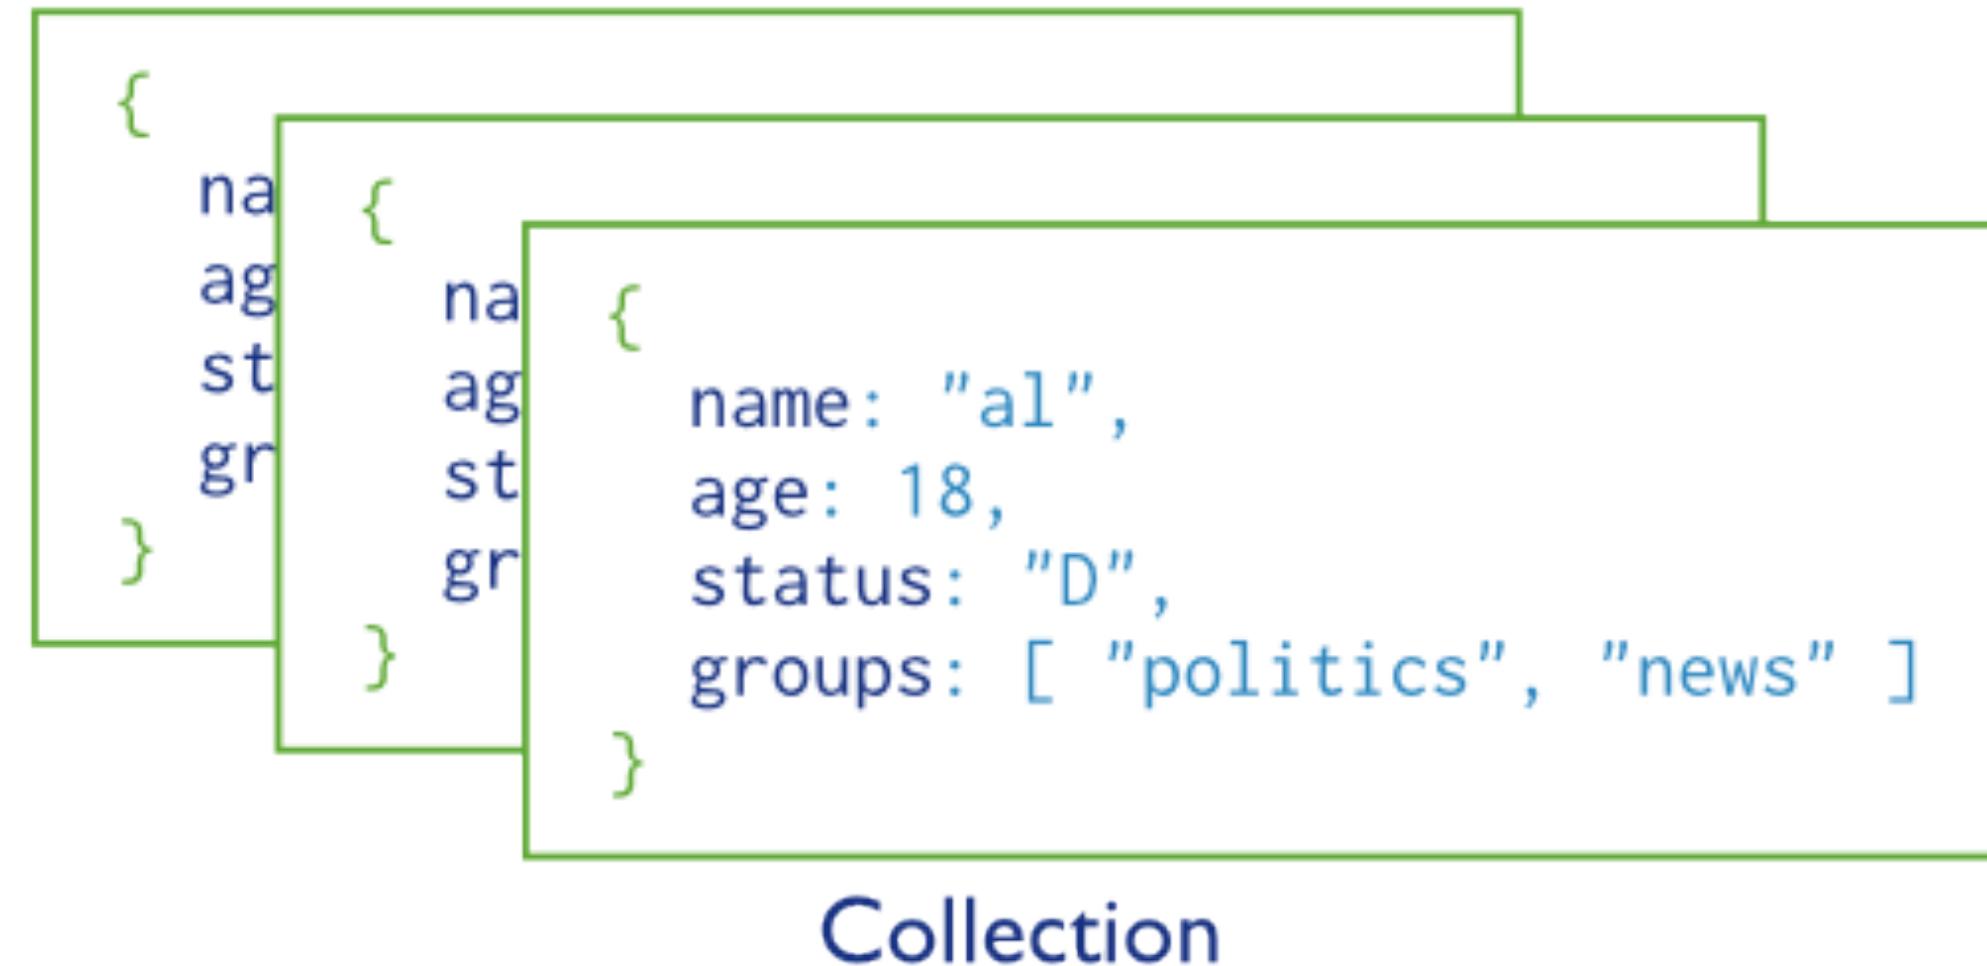
## Structure of a JSON-document:

```
{  
    name: "sue",  
    age: 26,  
    status: "A",  
    groups: [ "news", "sports" ]  
}  
  
← field:value  
← field:value  
← field:value  
← field:value
```

The value of field:

- Native data types
- Arrays
- Other documents

# Collections of Documents



Rule: Every document must have an `_id`.

# Embedded documents:

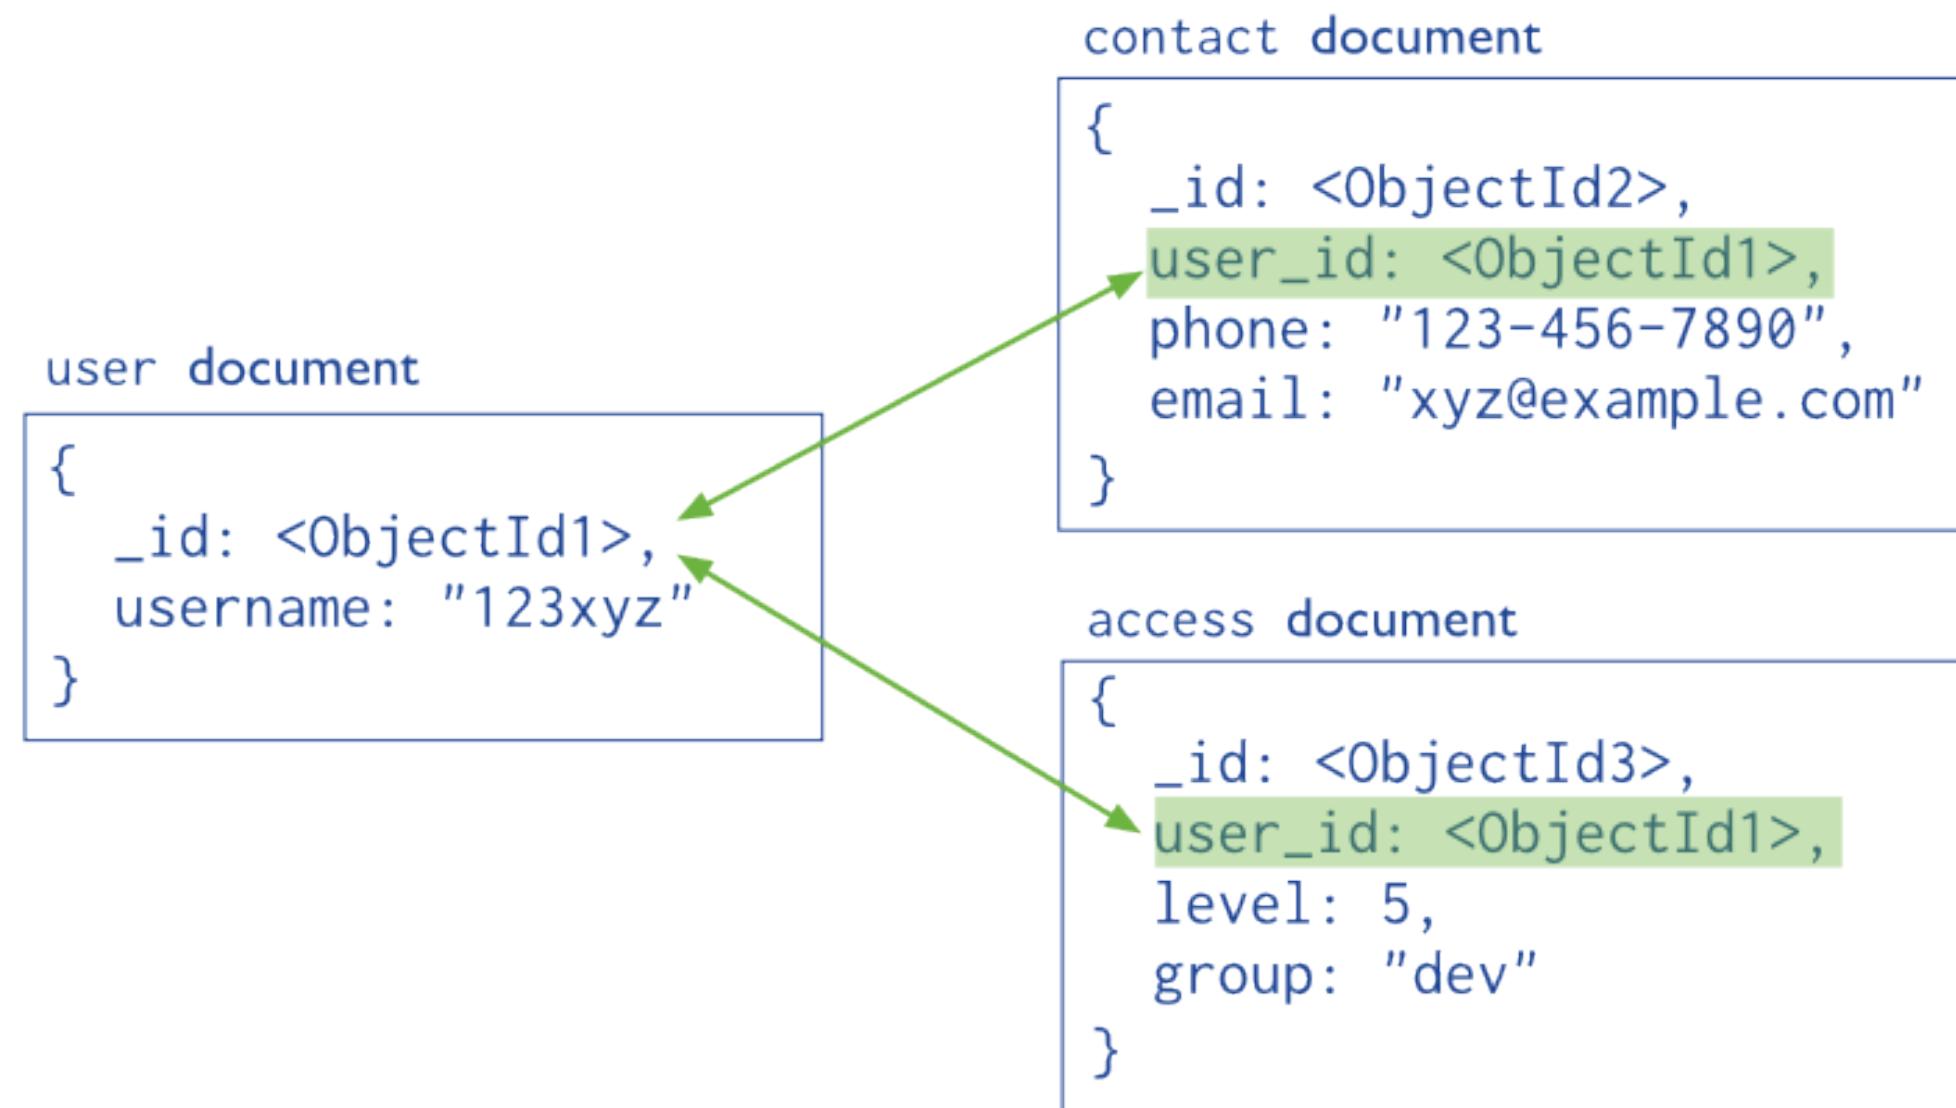
```
{  
  _id: <ObjectId1>,  
  username: "123xyz",  
  contact: {  
    phone: "123-456-7890",  
    email: "xyz@example.com"  
  },  
  access: {  
    level: 5,  
    group: "dev"  
  }  
}
```



Embedded sub-document

Embedded sub-document

# Reference documents



## Storage: BSON Format

- Binary-encoded serialization of JSON-like documents optimized for space and speed
- BSON types are a superset of JSON types<sup>94</sup>
- Zero or more key/value pairs are stored as a single entity<sup>93</sup>
- Large entities are prefixed with a length field to facilitate scanning

### BSON:

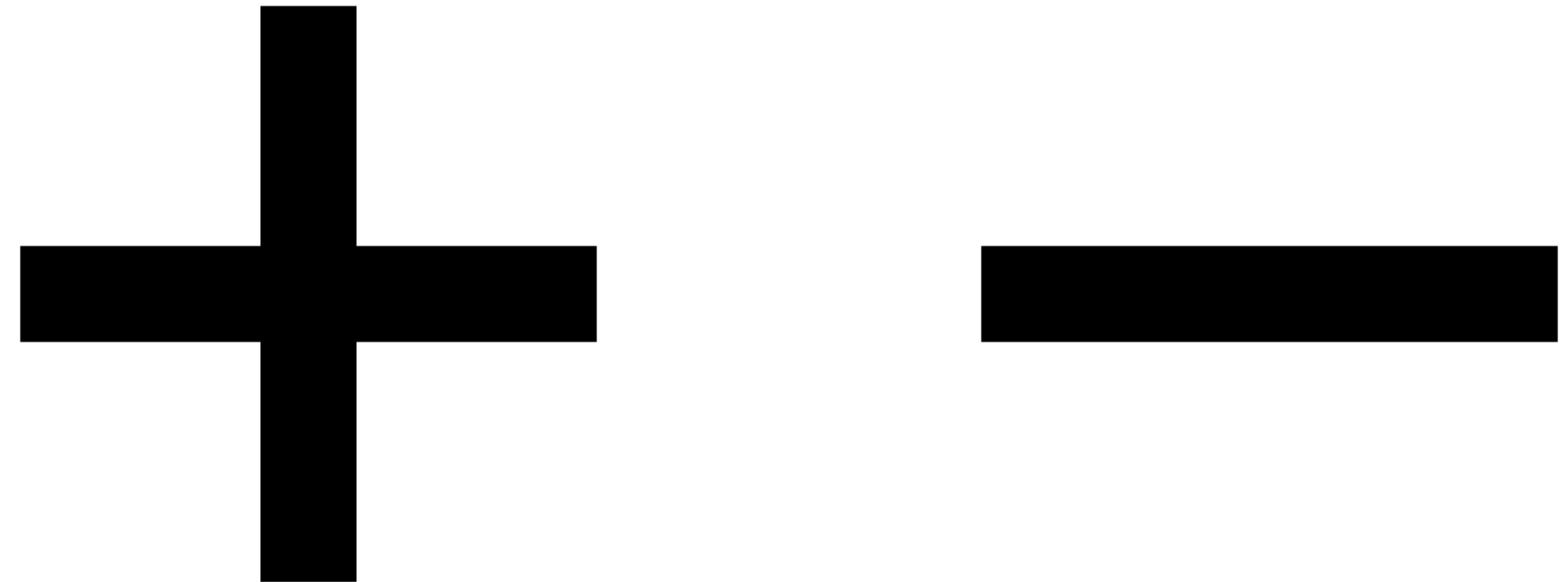
```
\x16\x00\x00\x00  
 \x02  
  name\x00  
  \x06\x00\x00\x00Devang\x00  
  \x00
```

```
// total document size  
// 0x02 = type String  
// field name  
// field value
```

---

<sup>94</sup> JSON does not have a date or a byte array type, for example

<sup>93</sup> Each entry consists of a field name, a data type, and a value



# Operations



# Create

Create a database

```
use database_name
```

Create a collection

```
db.createCollection(name, options)
```

# Insert

```
db.<collection_name>.insert({"name": "nguyen", "age": 24, "gender": "male"})
```

```
db.employee.insert({  
    name: "sally",  
    salary: 15000,  
    designation: "MTS",  
    teams: ["cluster-management"]  
})`
```

# Read

```
db.<collection_name>.find().pretty()  
  
db.employee #collection  
  .find( {  
    salary: {$gt:18000}, #condition  
    {name:1} #projection  
  })  
  .sort({salary:1}) #modifier
```

# Update

```
db.employee #collection
.update(
  {salary:{$gt:18000}}, #Update Criteria
  {$set: {designation: "Manager"}}, #Update Action
  {multi: true} #Update Option
)
```

## Delete

```
db.employee.remove(  
    {salary:{$lt:10000}}, #Remove Criteria  
)
```

# Aggregates

SQL-like aggregation functionality

Pipeline documents from a collection pass through an aggregation pipeline

Expressions produce output documents based on calculations performed on input documents

Example:

```
db.parts.aggregate(  
  {$group : {_id: type, totalquantity :  
    { $sum: quantity}  
  }})
```

## Save

```
db.employee.save(  
  { _id:ObjectId('string_id'),  
    "name": "ben",  
    "age": 23,  
    "gender":  
    "male"  
)
```

## Drop

- Drop a database
- Drop it:  
db.dropDatabase()
- Drop a collection:

db.<collection\_name>.drop()

# Mapping to SQL

## SQL Statement

SELECT \* FROM table

## MongoDB commands

db.collection.find()

SELECT \* FROM table WHERE artist = 'Nirvana' db.collection.find({Artist:"Nirvana"})

SELECT\* FROM table ORDER BY Title

db.collection.find().sort>Title:1

DISTINCT

.distinct()

GROUP BY

.group()

>=, <

\$gte, \$lt

# Comparison Operators

Name	Description
\$eq	Matches value that are equal to a specified value
\$gt, \$gte	Matches values that are greater than (or equal to) a specified value
\$lt, \$lte	Matches values less than or (equal to) a specified value
\$ne	Matches values that are not equal to a specified value
\$in	Matches any of the values specified in an array
\$nin	Matches none of the values specified in an array
\$or	Joins query clauses with a logical OR returns all
\$and	Join query clauses with a logical AND
\$not	Inverts the effect of a query expression
\$nor	Join query clauses with a logical NOR
\$exists	Matches documents that have a specified field

[source](#)

# Indexes

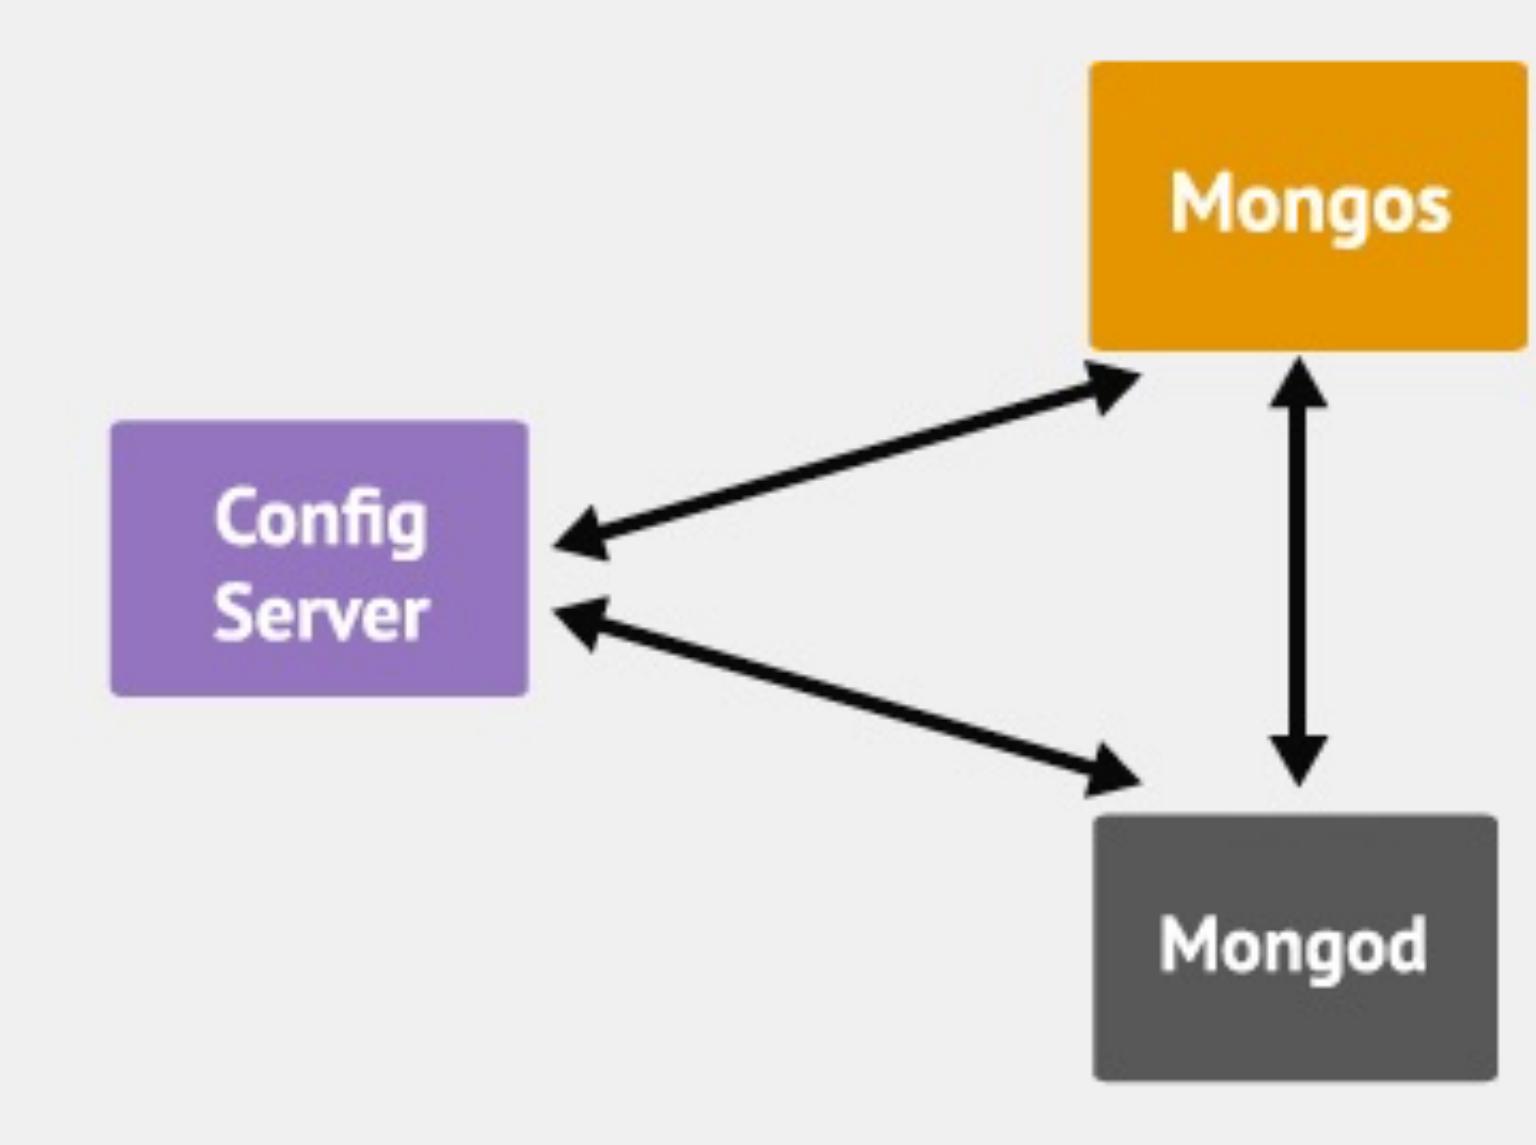
- B+ tree indexes
- An index is automatically created on the `_id` field (the primary key)
- Users can create other indexes to improve query performance or to enforce Unique values for a particular field
- Supports single field index as well as Compound index
- Like SQL order of the fields in a compound index matters
- If you index a field that holds an array value, MongoDB creates separate index entries for every element of the array

# Sparse Indexes

- Sparse\_- property of an index ensures that the index only contain entries for documents that have the indexed field. (so ignore records that do not have the field defined)
- If an index is both unique and sparse – then the system will reject records that have a duplicate key value but allow records that do not have the indexed field defined

# Architecture

- Mongod – Database instance
- Mongos - Sharding processes:
  - Like a database router processes all requests
  - Decides how many and which *mongod* should receive the query
  - No local data
  - Collects the results, and sends it back to the client.
- Config Server
  - Stores cluster chunk ranges and locations
  - Can have only 1 or 3 (production must have 3)



# Mongod and Mongos



# Client

- Mongo – an interactive shell ( a client)
- Fully functional JavaScript environment for use with a MongoDB
- You can have one mongos for the whole system no matter how many mongods you have
- OR you can have one local mongos for every client if you wanted to minimize network latency.

# Replication

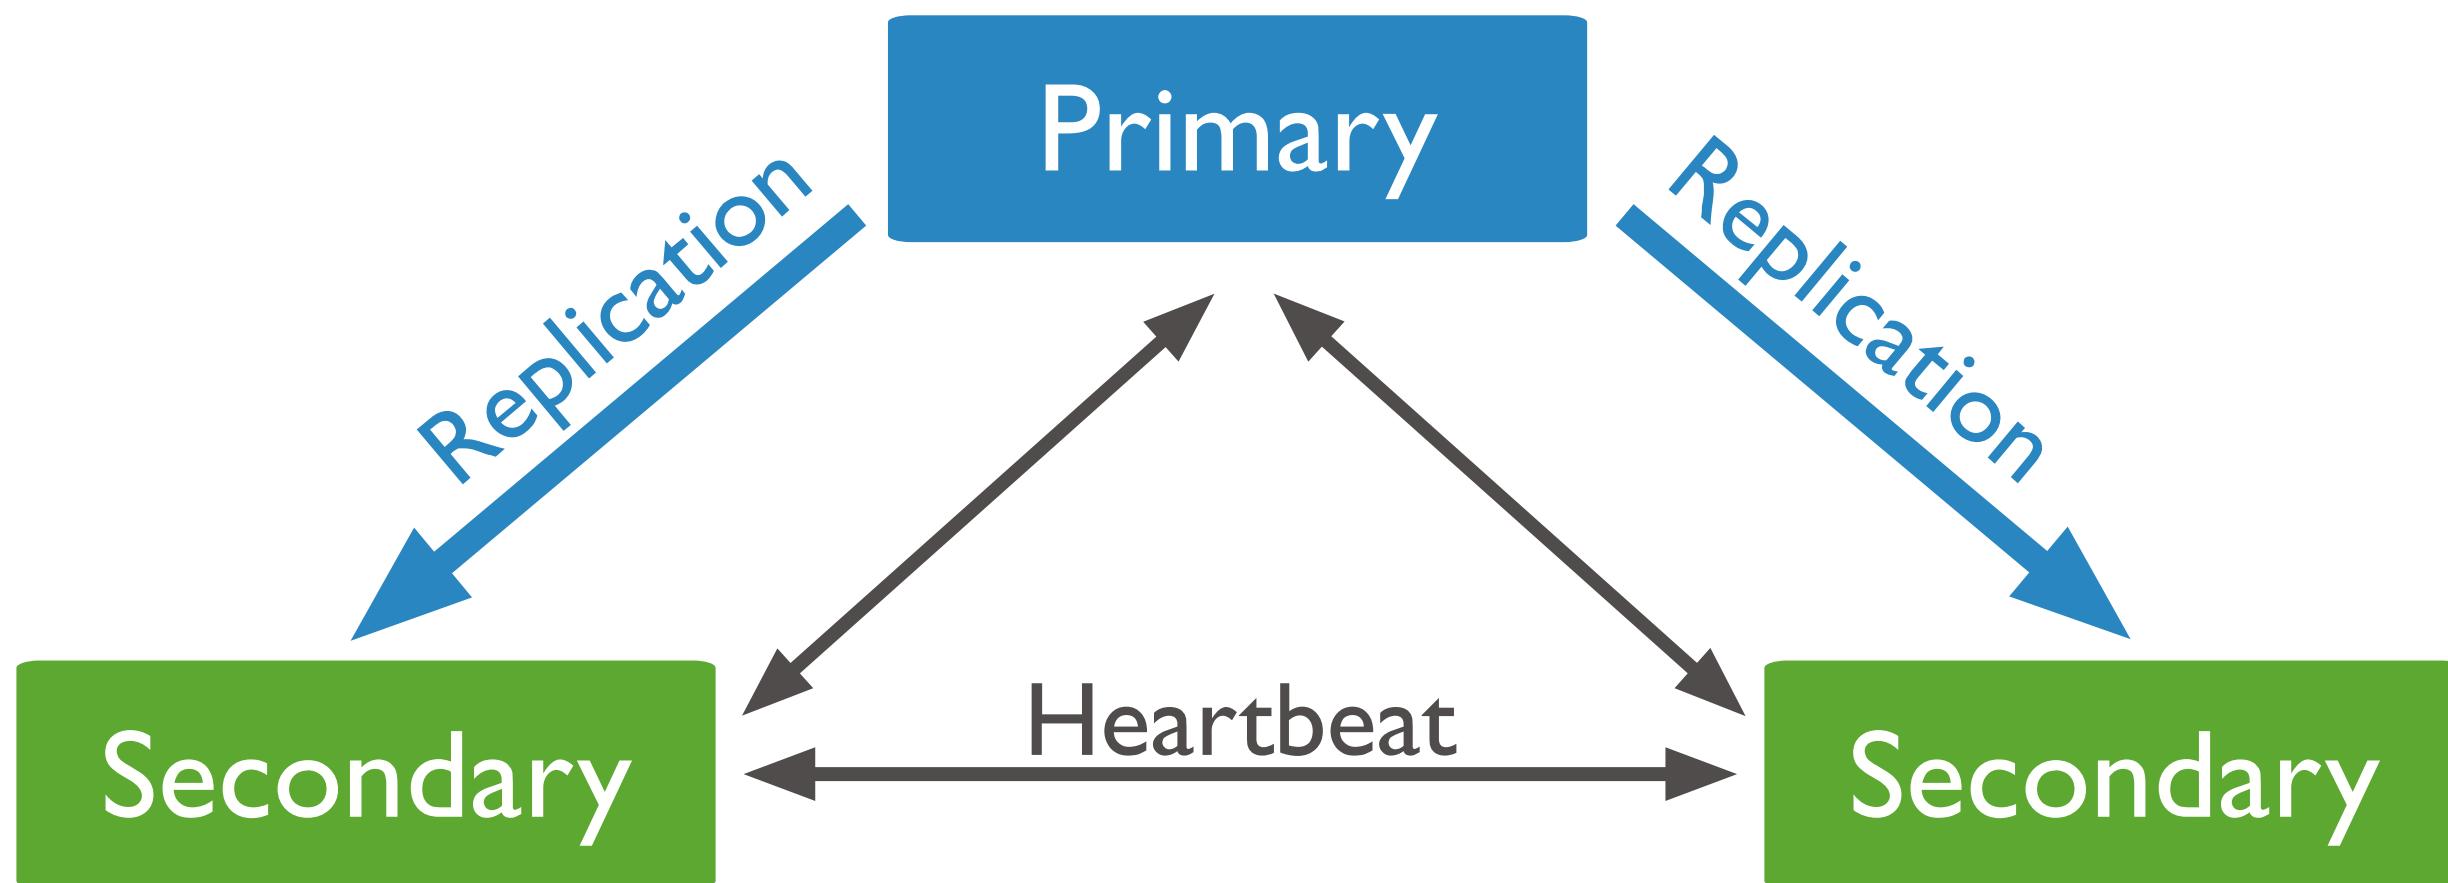
For redundancy MongoDB provides asynchronous replication.

Only one database node is in charge of write operations at any given time (called primary server/node).

Read operations may go to this same server for strong consistency semantics or to any of its replica peers if eventual consistency is sufficient.

# Master Slave Replication

Consists of two servers out of one which takes the role of a master handling write requests and replicating those operations to the second server, the slave.



# Replica Sets

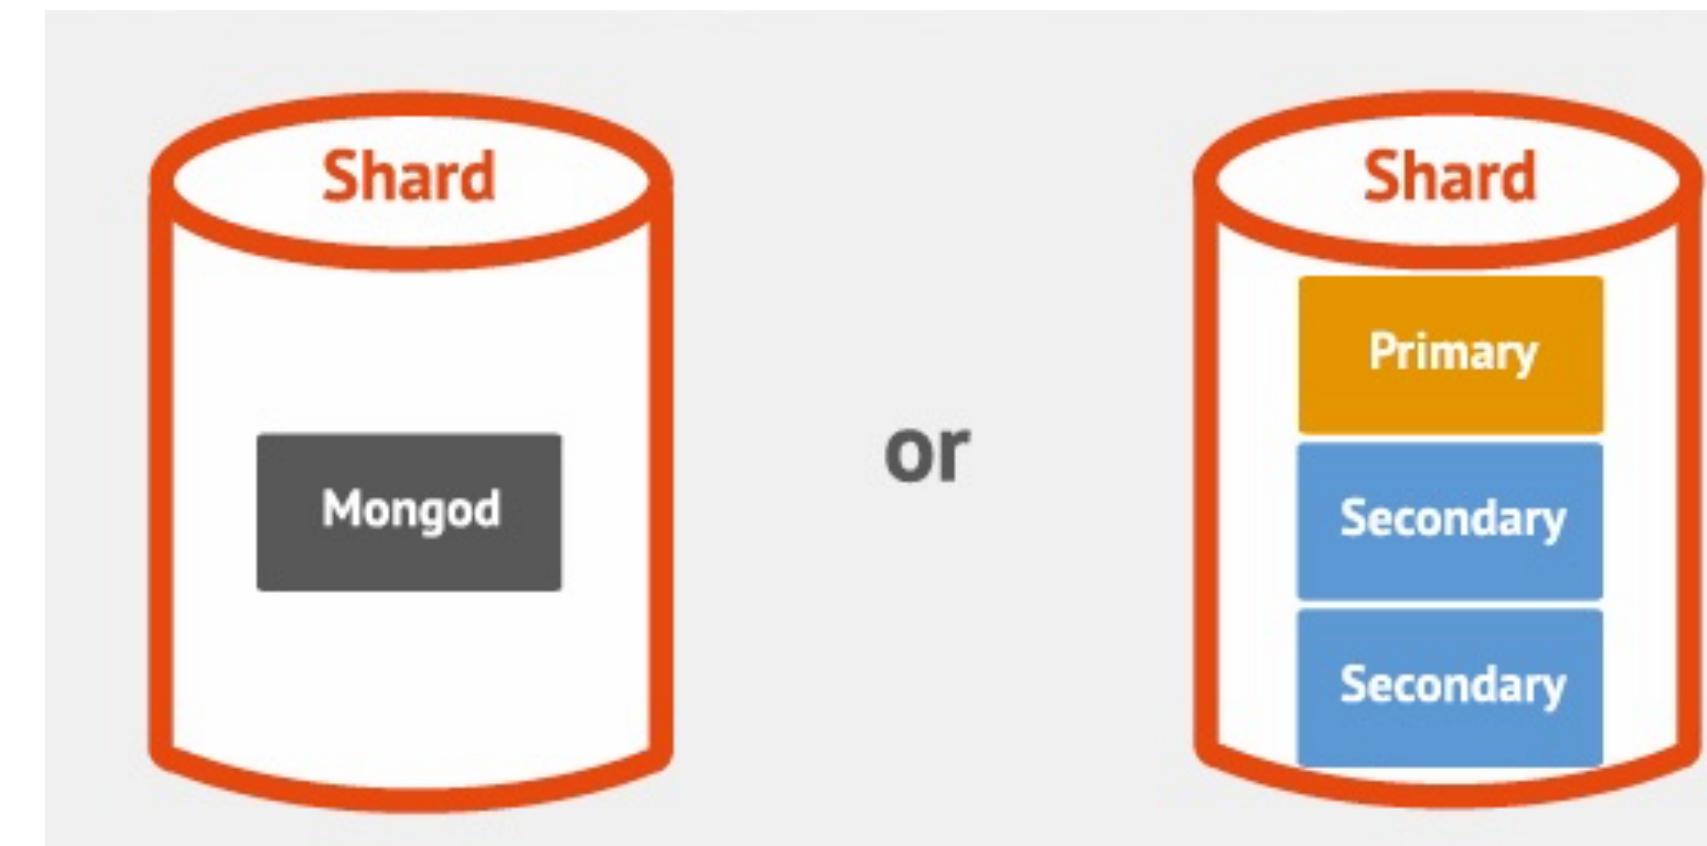
Consists of groups of MongoDB nodes that work together to provide automated failover.

# Partitioning

- called Sharding in MongoDB
- User defines shard key for partitioning
- Shard key defines range of data
- Key space is like points on a line
- Range is a segment of that line

## What is a Shard?

- Shard is a node of the cluster
- Shard can be a single mongod or a replica set
- Default max chunk size: 64mb
- MongoDB automatically splits & migrates chunks when max reached



# Auto-sharding

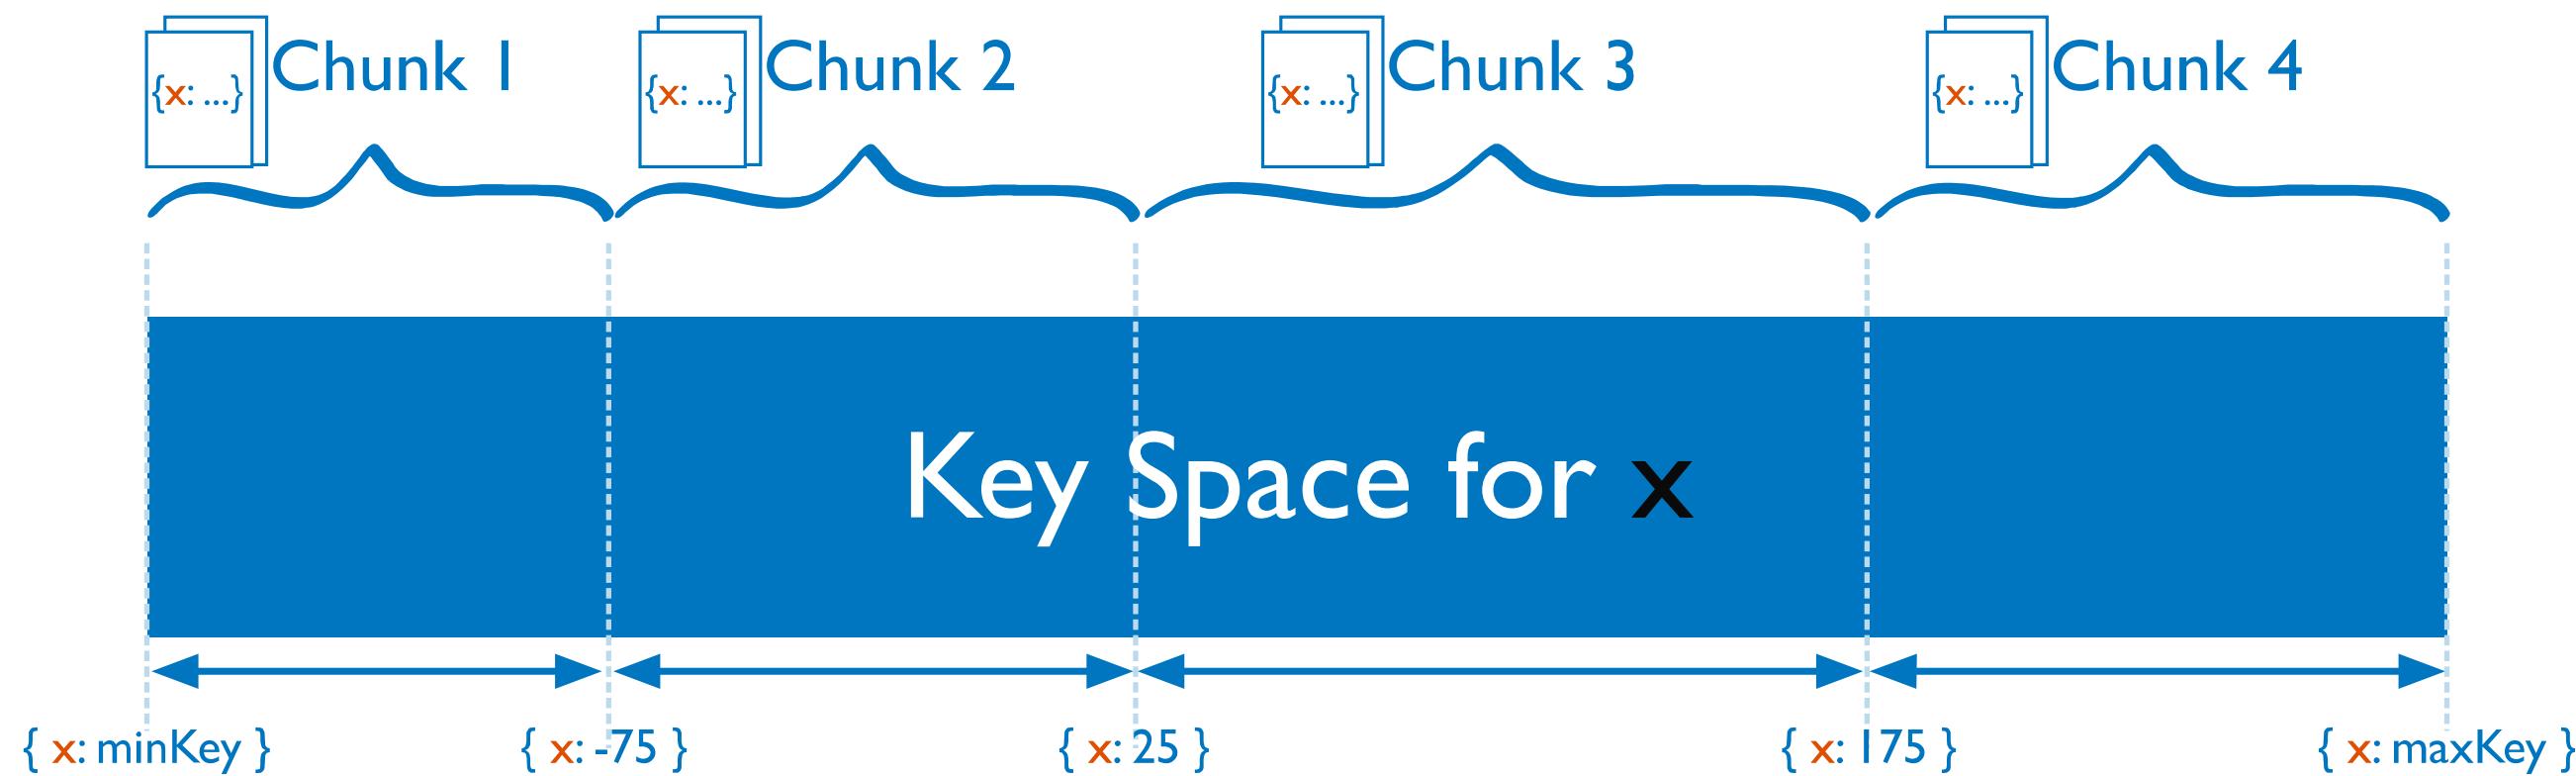
- Minimal effort required
  - Enable Sharding for a database
  - Shard collection within database
  - Decide Sharding Strategy

# MongoDB Sharding Strategies

- Ranged
- Hashed
- Tag-aware

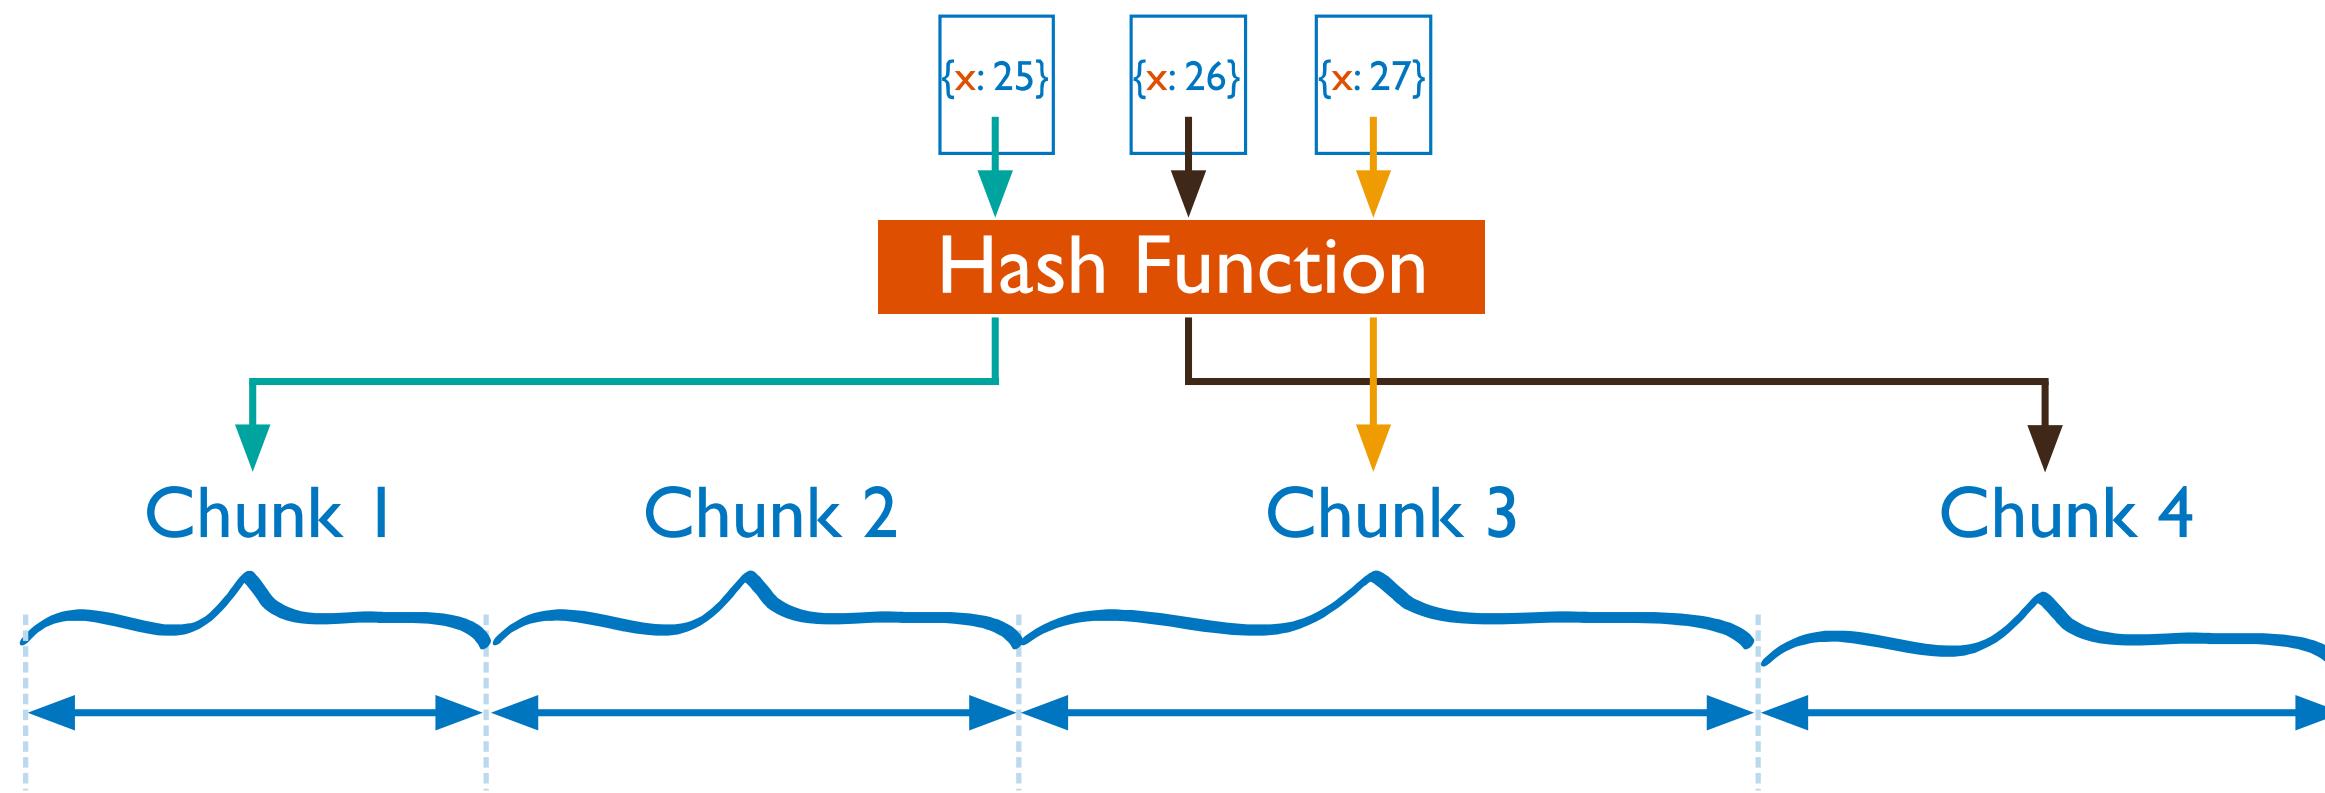
## Range Sharding

- Splits shards based on sub-range of a key (or also multiple keys combined)
  - Simple Shard Key: {deviceId}
  - Composite Shard Key: {deviceId, timestamp}



## Hash Sharding

- MongoDB applies a MD5 hash on the key when a hash shard key is used:
  - Hash Shard Key(deviceId) = MD5(deviceId)
  - Ensures data is distributed randomly within the range of MD5 values



## Tag Sharding

Tag-aware sharding allows subset of shards to be tagged, and assigned to a sub-range of the shard-key.

Example: Sharding User Data belong to users from 100 “regions”

Collection: Users, Shard Key: {uld, regionCode}

Tag based on macro regions

## Tag Sharding Example

**Collection: Users, Shard Key: {uld, regionCode}**

Tag	Start	End
West	MinKey, MinKey	MaxKey, 50
East	MinKey, 50	MaxKey, MaxKey
Shard1, Tag=West	Shard2, Tag=West	Shard3, Tag=East
Shard4, Tag=East		

The diagram illustrates the distribution of data regions across four shards. Each shard contains a Primary node and two Secondary nodes. The shard ranges are defined by the Start and End values in the table above. Regions 1-50 are assigned to the West, and regions 51-100 are assigned to the East.

Assign Regions  
1-50 to the West

Assign Regions  
51-100 to the  
East

## Which Sharding to use?

### Usage

Scale

Geo-Locality

Hardware Optimization

Lower Recovery Times

### Required Strategy

Range or Hash

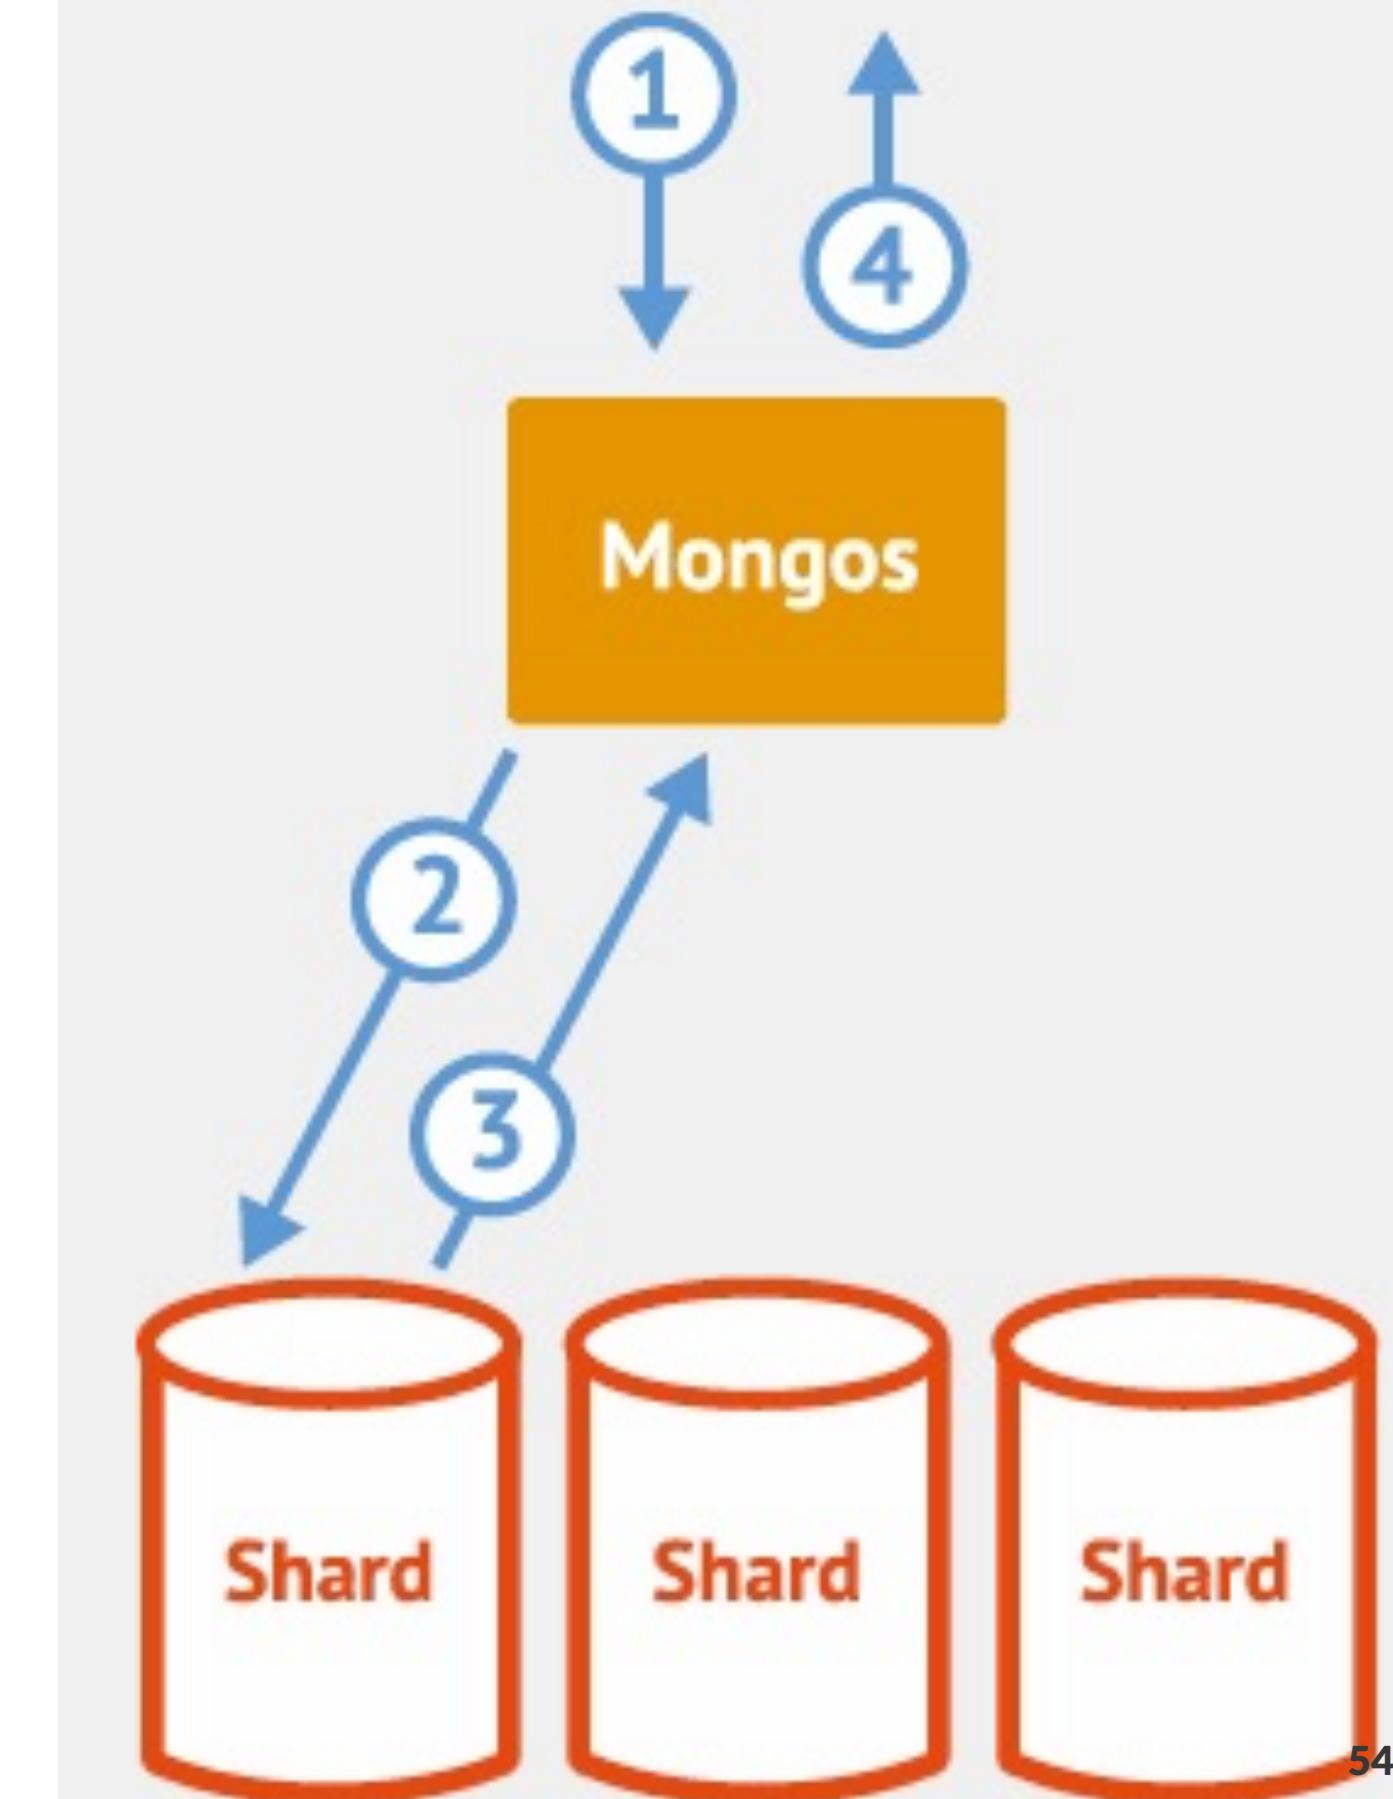
Tag-aware

Tag-aware

Range or Hash

## Routing and Balancing

- Queries routed to specific shards
- MongoDB balances cluster
- MongoDB migrates data to new nodes



# MongoDB Security

- SSL
  - between client and server
  - Intra-cluster communication
- Authorization at the database level
  - Read Only/Read+Write/Administrator



## References

- [Mongodb.com](#)
- No SQL Distilled by P. Sadalage and M. Fowler
- MongoDB Applied Design Patterns by R. Copeland
- The Definitive Guide to MongoDB by Plugge, Membry and Hawkins

